

Hi, I'm Nathan

I go by @myobie online

on almost every network

There will be Q&A

I design distributed
applications from UI to
data storage

I built this presentation
using choo and phoenix

choo.io

This talk is interactive

If you have cellular data
then maybe let others
use the wifi

elixir.myobie.com

You can rate how you feel
about this talk anytime
using the slider



The average rating will
frequently update and
will appear on top of the
slides



You can also submit an emoji that will show up in the bottom right corner



Have fun

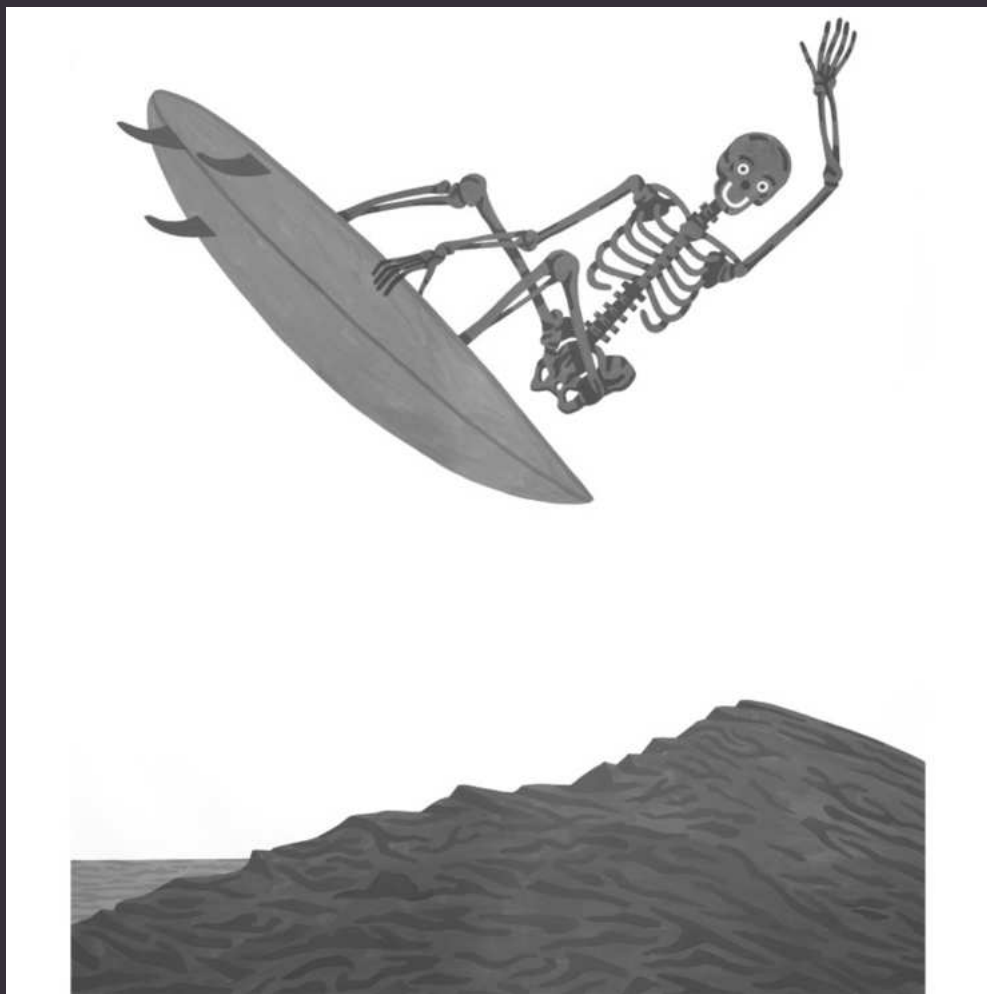
Distributed applications
are hard

Network partitions

Concurrency

Locking (Linearization)

Stale data



by Ted Parker

<http://ted-parker.com/portfolio/surfs-dead/>

Network latency can
cause stale reads from
consistent systems

GET

POST

NEW

OLD



OLD

NEW

*Happens
all the time*



Scaling is hard

Load balancing

Parallelism

Databases

AUTOVACUUM & VACUUM

*Managing databases
is still hard*



Why build a distributed
system in the first place?

Horizontal scaling

Support “billions” of users

Offline

Bad networks



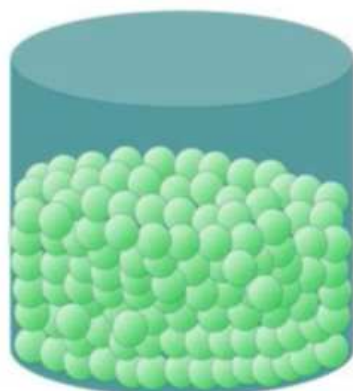
*I watched a talk by
Joe Armstrong recently*

<https://youtu.be/bo5WL5IQAd0>

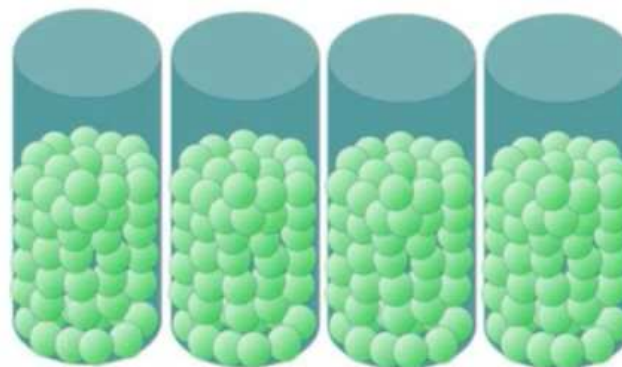


But packing sand into boxes is easy

Source
Erlang factory 2010
San Francisco
Patrik Nyblom



Single core



Multi core

granular = scalable





Crazy Idea

Having a central database
makes some things easier

However, a single
database is hard to scale

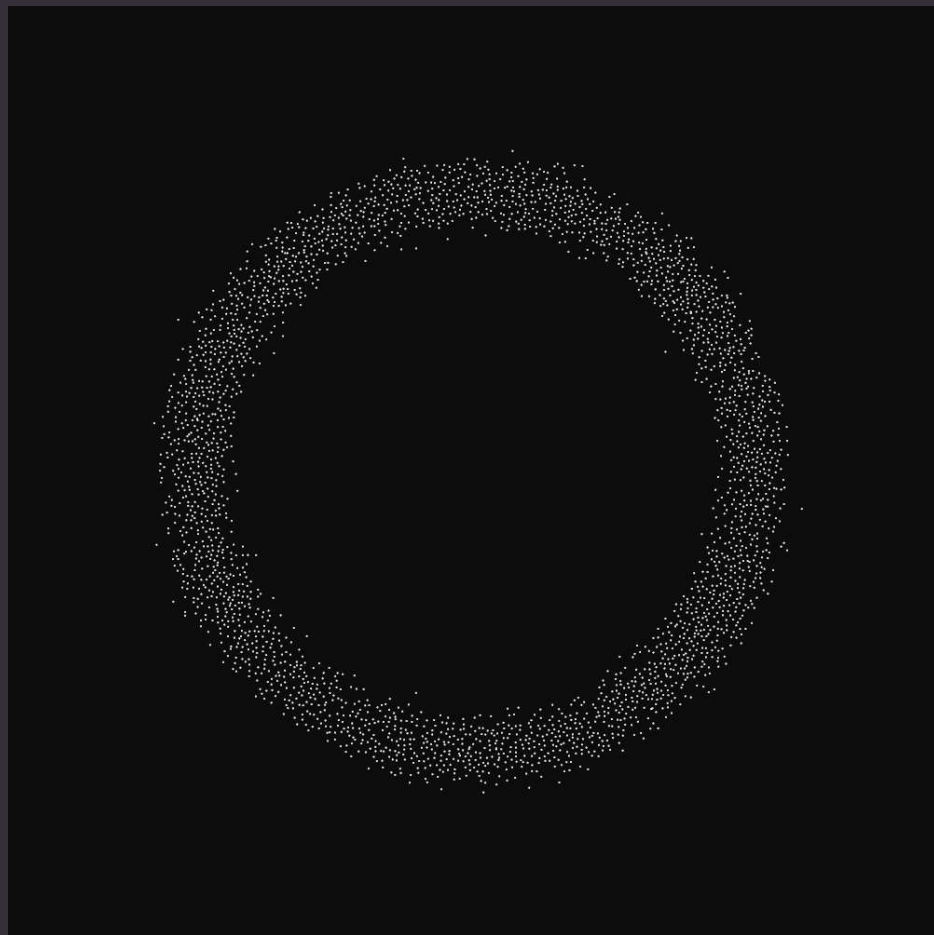
What if every user had
their own database?
Stored on S3 as a file?!



by Charis Tsevis
<https://flic.kr/p/8hUVRV>

Time

Human time cannot be
trusted in a distributed
system



by Anders Hoff

<https://twitter.com/inconvergent>

*Any working distributed
system is an implementation
of how that organization has
understood time*

Actors



by Nagano Toyokazu
<https://www.flickr.com/photos/toyokazu>

At Wunderlist
we used actors to wrap
each websocket
connection

A great way to
contain state and
linearize access

Wait a second?!



Elixir processes

What if every user had
their own process?!

To trust a process as the
authority for a user ...

We need to know there
will be only one per user

*If we were to accidentally
have two “myobie” processes,
then we have a split brain*

:global

Not consistent

Are you sure?

If any name clashes are discovered, function `Resolve` is called.

<http://erlang.org/doc/man/global.html>



: pg2

Not consistent

Are you sure?

*pg2 replicates all name
lookup information in a
way that doesn't require
consistency ...*

<http://erlang.org/pipermail/erlang-questions/2012-June/067220.html>



:gproc

Same

Are you sure?

*While gproc has been tested
very thoroughly ... its
reliance on gen_leader is
problematic.*

[https://christophermeiklejohn.com/erlang/2013/06/05/
erlang-gproc-failure-semantics.html](https://christophermeiklejohn.com/erlang/2013/06/05/erlang-gproc-failure-semantics.html)

Does gen_leader
have problems?

I don't know.



Elixir.Registry

Single vm

Consistent

*It's “easy” to be consistent in
one process on one vm*



If that one vm stops ...

... then our entire registry
is gone 💥



How do we make 100%
certain we never
accidentally boot two
vms?



Zookeeper

Consistent

Pretty available

Partition tolerant

Also: includes recipes

<https://zookeeper.apache.org/doc/current/recipes.html>

Recipe: distributed locks

Major





by Meghan Roberts

<https://www.flickr.com/photos/88009602@N05/8696887207>

Battle-tested system for
getting time under
control

We can know if one thing
happens before another

Highlander

github.com/myobie/highlander

There can be only one

90's TV show reference



“Objects” persist their
state to S3 during update

“Objects” read from S3
during init

“Objects” respond to calls
using their in-memory
state like GenServers

“Objects” teardown after
inactivity

Process registry stored in
zookeeper

GenServer *:via*

```
GenServer.call(  
  {:via, MyRegistry, {:user, user_id}},  
  :do_stuff  
)
```

```
defmodule MyRegistry do
  def send(name, message) do
    end

  def whereis_name(name) do
    end

  def register_name(name) do
    end

  def unregister_name(name) do
    end
end
```

Zookeeper is a simple key value store. The key will be the object's id and the value the node name.

```
defmodule MyRegistry do
  def whereis_name(name) do
    case ZK.get_node_name(name) do
      # ...
    end
  end
end
```

```
defmodule MyRegistry do
  def register_name(name) do
    case ZK.create_znode(name) do
      # ...
    end
  end
end
```



```
defmodule ZK do
  def create_znode(name) do
    {:ok, pid} = ZNode.start_link(name)

    if ZNode.first?(pid) do
      {:ok, pid}
    else
      :ok = ZNode.delete(pid)
      {:error, :already_exists}
    end
  end
end
```

```
defmodule ZNode do
  def init(name) do
    {:ok, path} =
      Zookeeper.Client.create(:zk,
        prefix(name, UUID.uuid4(:hex)),
        to_string(Node.self),
        makepath: true,
        create_mode: :ephemeral_sequential)
    {:ok, %{path: path, name: name}}
  end
end
```

```
defmodule ZNode do
  def init(name) do
    {:ok, path} =
      Zookeeper.Client.create(:zk,
        prefix(name, UUID.uuid4(:hex)),
        to_string(Node.self),
        makepath: true,
        create_mode: :ephemeral_sequential)
    {:ok, %{path: path, name: name}}
  end
end
```

Zookeeper will keep time
in order

```
opts = [host,  
        [stop_on_disconnect: true, name: :zk]]  
  
children = [  
    worker(Zookeeper.Client, opts, []),  
    worker(Registry.Server, [], []),  
    worker(Registry.NodeCycleServer, [], []),  
    supervisor(Object.Supervisor, [], [])  
]  
  
supervise children, strategy: :rest_for_one
```

OTP will keep the
processes running

Macros

```
defmodule Todo do
  use Highlander.Object, type: :todo

  defobject title: "",
    completed: false,
    color: :blue
end
```



```
id = UUID.uuid4()
```

```
{:ok, todo} = Todo.get id
```

```
assert todo.completed == false  
assert todo.title == ""  
assert todo.color == :blue
```

```
todo = %{todo | title: "Hello"}
```

```
{:ok, _} = Todo.put id, todo
```

```
# any amount of time later
```

```
{:ok, todo} = Todo.get id
```

```
assert todo.title == "Hello"
```

Model everything as a
process

```
defmodule User do
  use Highlander.Object, type: :user

  defobject name: "",
    email: ""
end
```

What about shared
objects?

```
defmodule List do
  use Highlander.Object, type: :list

  defobject title: "",
    todos: []
end
```

```
defmodule Todo do
  defstruct title: "",
    completed: false,
    color: :blue
end
```


Like an app where you
can share a todo list with
other people...

User ↔ List ↔ User

```
defmodule User do
  use Highlander.Object, type: :user

  defobject name: "",
    email: "",
    known_lists: []
end
```

```
defmodule List do
  use Highlander.Object, type: :list

  defobject title: "",
    todos: []
    allowed_users: []
end
```

```
defmodule List do
  def handle_update(:insert_todo, list,
                    todo, user) do
    if Enum.any?(list.allowed_users,
                  fn id -> id == user.id end) do
      %{todos: [todo | state.todos]}
    end
  end
end
```

One list's activity doesn't
slow down another

Aside: multi-vm testing

I copied how
phoenix does it


```
def spawn do
  :net_kernel.start([:"primary@127.0.0.1"])
  :erl_boot_server.start([])
  allow_boot to_char_list("127.0.0.1")

  Application.get_env(:highlander,
    :spawn_nodes, [])
  /> Enum.map(&Task.async(fn ->
    {:ok, _node} = spawn_node(&1)
  end))
  /> Enum.map(&Task.await(&1, 30_000))
end
```

```
defp spawn_node(node_host) do
  {:ok, node} = :slave.start(
    to_char_list("127.0.0.1"),
    node_name(node_host),
    inet_loader_args())

  add_code_paths(node)
  transfer_configuration(node)
  ensure_applications_started(node)
  {:ok, node}
end
```

test/support/cluster.ex

```
$ epmd -daemon
```



Scaling

Like pouring sand...

How many locks can a
small zookeeper cluster
handle?

I don't know.

Sharded zk

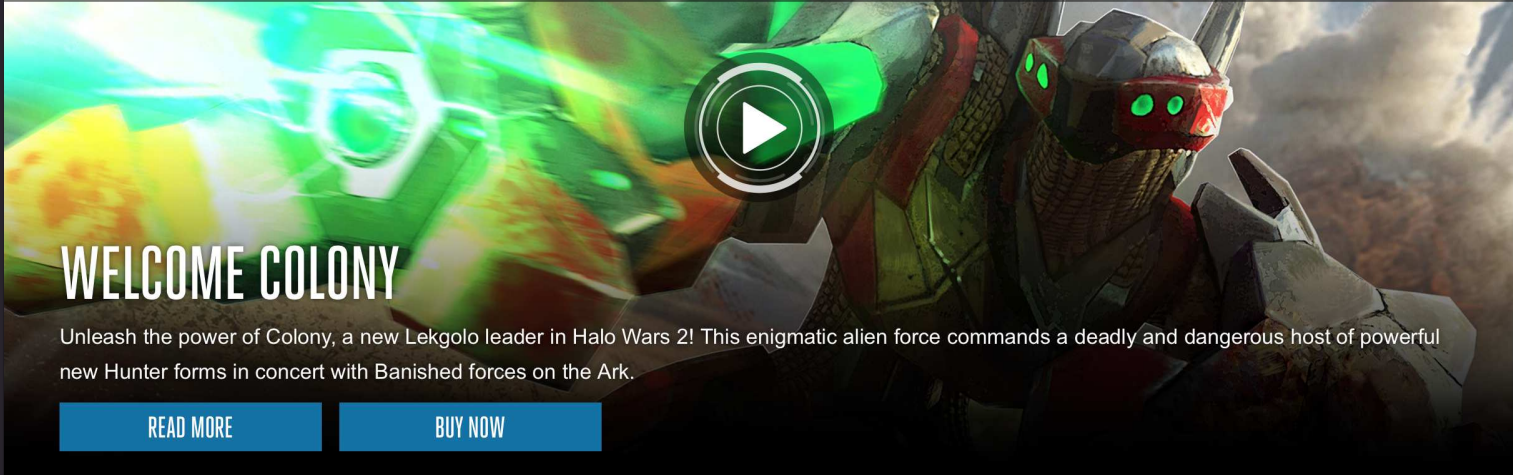
Sharded registry
keyspaces

Halo

halowaypoint.com

HALO GAMES NEWS ESPORTS COMMUNITY FORUMS UNIVERSE SHOP






Sign In




WELCOME COLONY


Unleash the power of Colony, a new Lekgolo leader in Halo Wars 2! This enigmatic alien force commands a deadly and dangerous host of powerful new Hunter forms in concert with Banished forces on the Ark.

READ MOREBUY NOW






HALO LEGENDS NOW AVAILABLE ON NETFLIX



HALO CRATE #6 FIGURE REVEALED



HCS COMING TO DAYTONA BEACH

Orleans

<https://dotnet.github.io/orleans/Documentation/Introduction.html>

MSFT Service Fabric

Uses zk to setup
partitions

Partitions actors and
services based on the zk
registry

Akka

Sure.

riak_core

Sure.



Have fun.

Have crazy ideas.

Build things.



Please ask any question,
even if you think it's
dumb or novice or
complicated or ...

Don't make a long
statement that isn't a
question.

Q&A