

# The Scalable Stateful Web

with Phoenix and Riak Core

Ben Tyler  
Booking.com

ElixirConf.EU 2016

What?

# What?

"joy of cool tech"

# What?

*"joy of cool tech"*

Talk notes and code will be available w/ the slides.

# What?

"joy of cool tech"

Talk notes and code will be available w/ the slides.

Let's chat after the conf!

# Buzzword Bingo

Stateful

# Buzzword Bingo

Stateful

Distributed

# Buzzword Bingo

Stateful

Distributed

Fault-tolerant



# Buzzword Bingo

Stateful

Distributed

Fault-tolerant

Real-time

# Buzzword Bingo

Stateful

Distributed

Fault-tolerant

Real-time

Impress your cat

# Buzzword Bingo

Stateful

Distributed

Fault-tolerant

Real-time

Impress your cat

(application)

# Buzzword Bingo

Stateful

Distributed

Fault-tolerant

Real-time

Impress your cat

(application)

# Stateful

Memory that lasts for more than one request

# Stateful

(stateless?)


Stateless

# Stateless







# Stateless



# Stateless




# Stateless




# Stateless

Horizontally  
scalable




# Stateless



Horizontally  
scalable

One stop shop

# Stateless

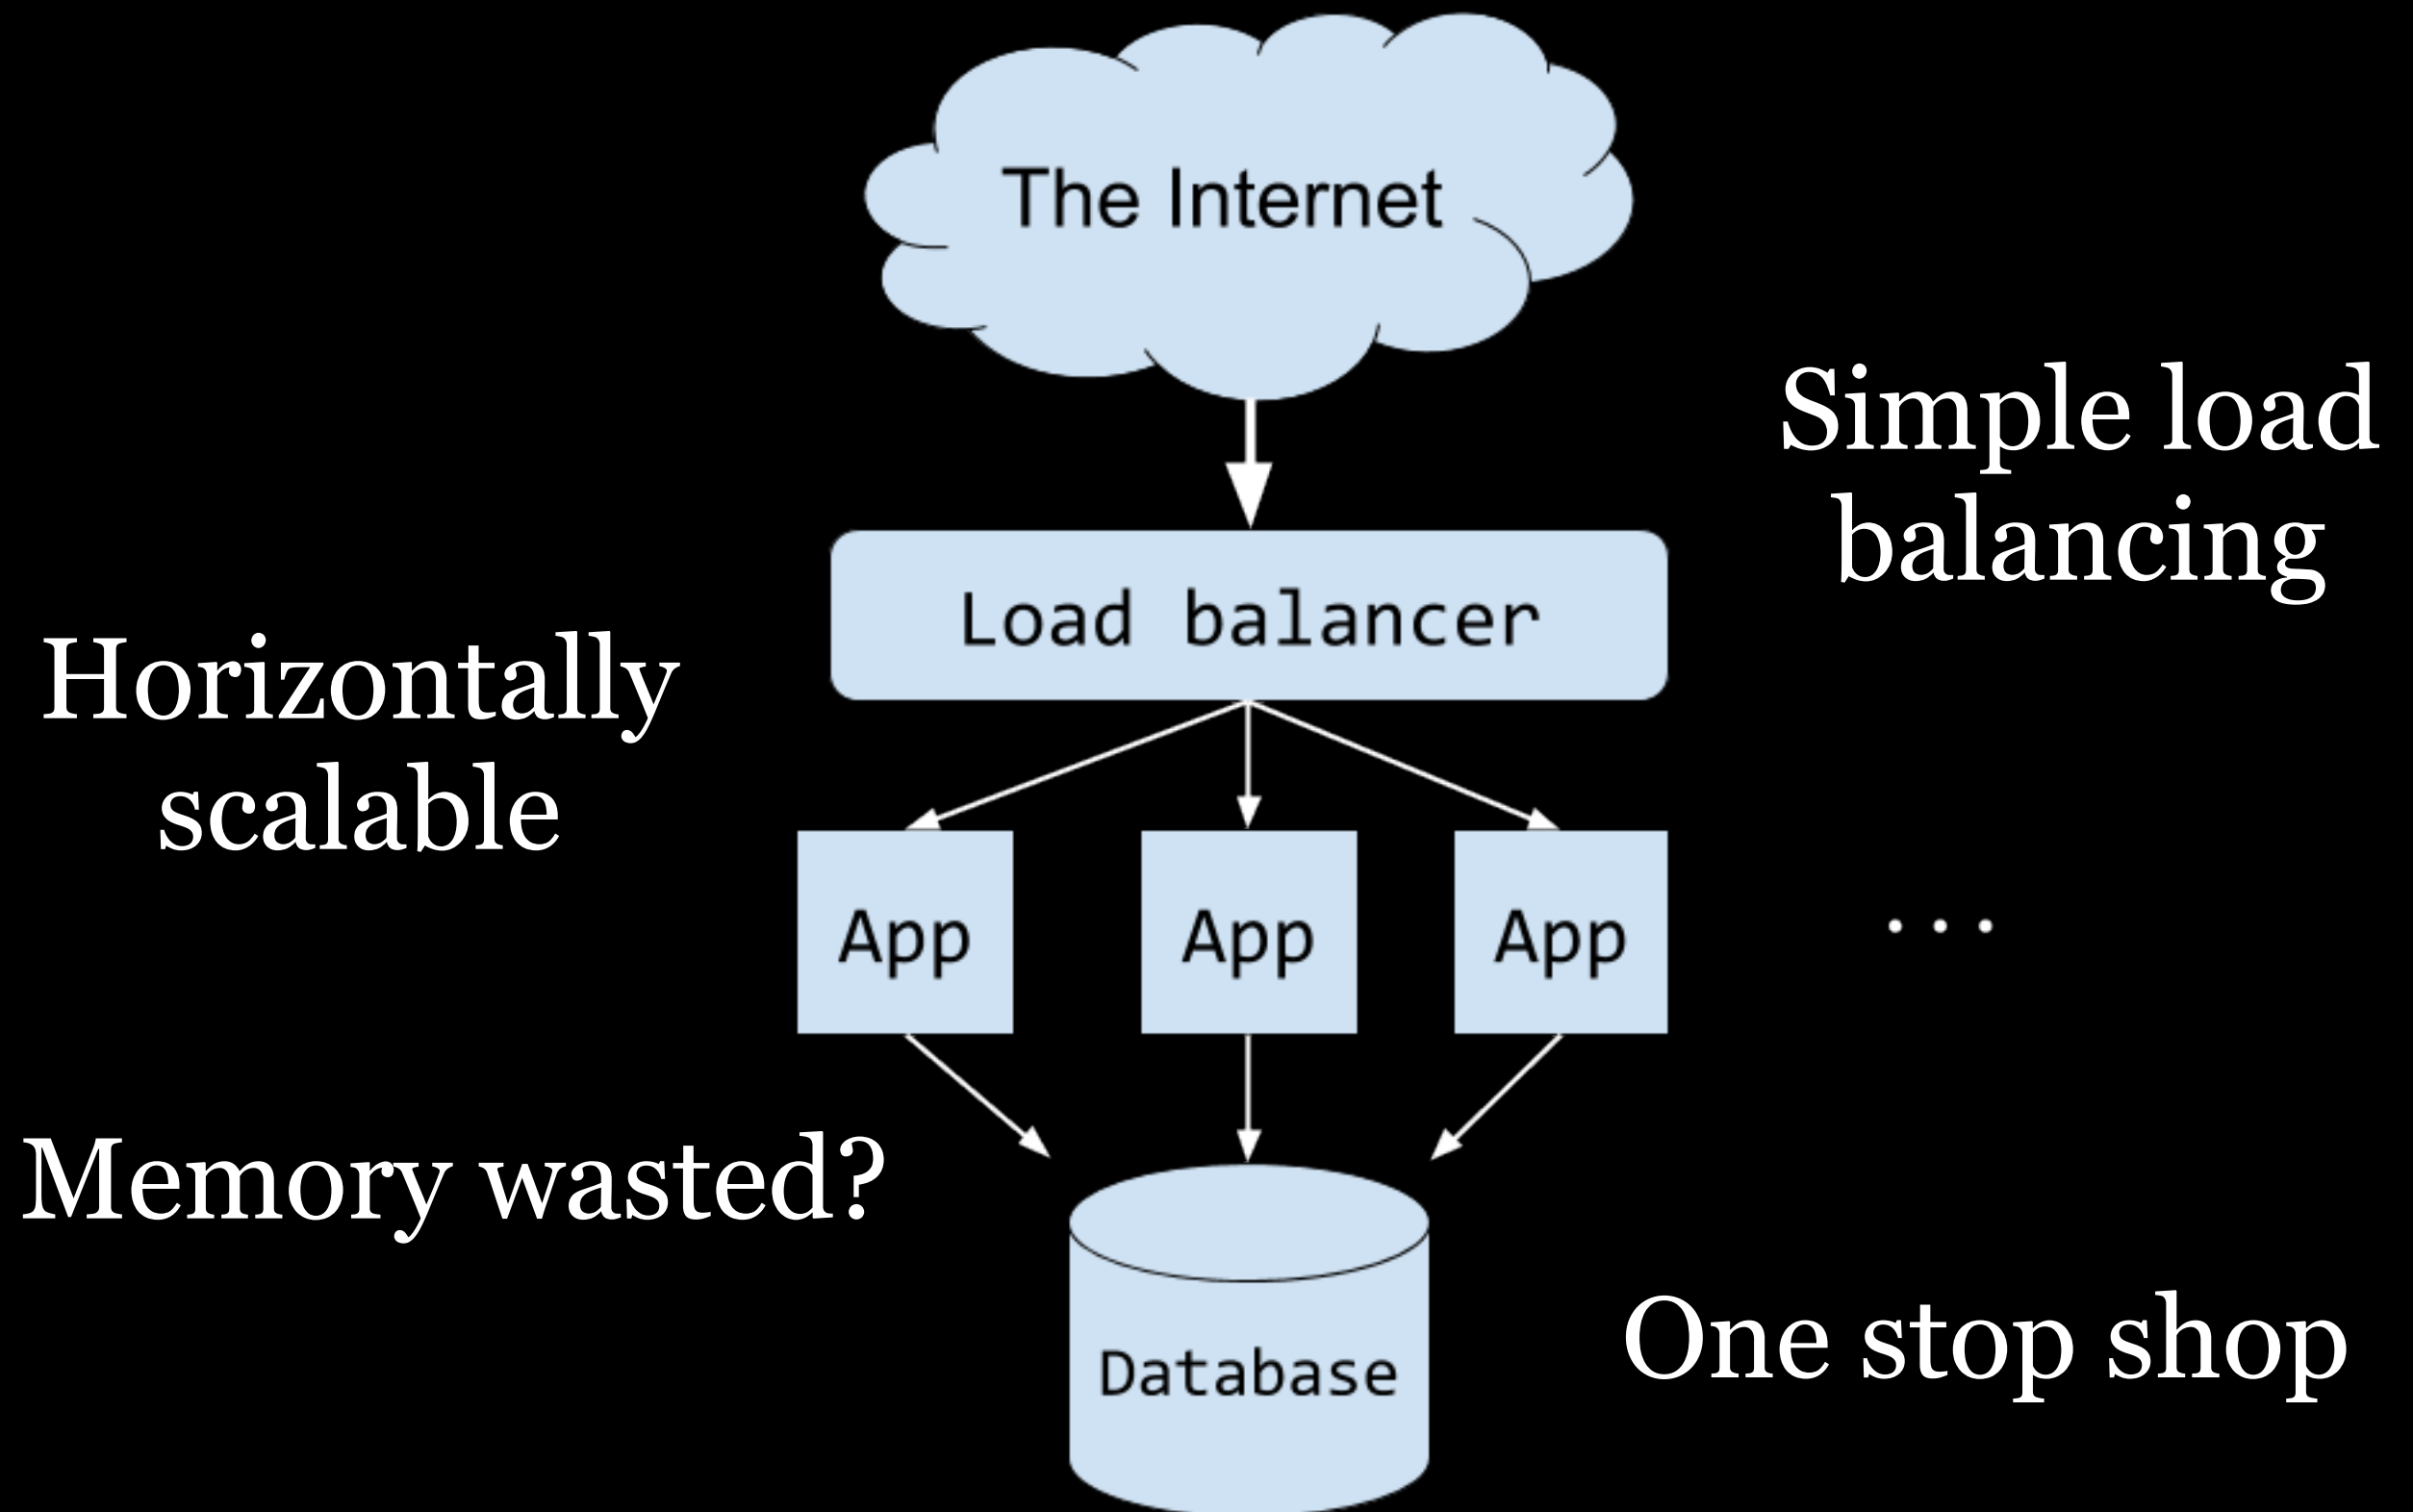


Horizontally  
scalable

Simple load  
balancing

One stop shop

# Stateless



# Latency Numbers

CPU L1 cache reference	0:00:01



# Latency Numbers

CPU L1 cache reference	0:00:01
Main memory reference	0:03:20

# Latency Numbers

CPU L1 cache reference	0:00:01
Main memory reference	0:03:20
Read 1MB sequentially from memory	6 days

# Latency Numbers

CPU L1 cache reference	0:00:01
Main memory reference	0:03:20
Read 1MB sequentially from memory	6 days
Network round trip, same datacenter	11.5 days

# Latency Numbers

CPU L1 cache reference	0:00:01
Main memory reference	0:03:20
Read 1MB sequentially from memory	6 days
Network round trip, same datacenter	11.5 days
Read 1MB sequentially from disk	463 days

# Latency Numbers

CPU L1 cache reference	0:00:01
Main memory reference	0:03:20
Read 1MB sequentially from memory	6 days
Network round trip, same datacenter	11.5 days
Read 1MB sequentially from disk	463 days

# Latency Numbers

CPU L1 cache reference	0:00:01
Main memory reference	0:03:20
Read 1MB sequentially from memory	6 days
Network round trip, same datacenter	11.5 days
Read 1MB sequentially from disk	463 days

# Latency Numbers


CPU L1 cache reference	0:00:01
Main memory reference	0:03:20
Read 1MB sequentially from memory	6 days <- do this!
Network round trip, same datacenter	11.5 days
Read 1MB sequentially from disk	463 days

# Latency Numbers


CPU L1 cache reference	0:00:01
Main memory reference	0:03:20
Read 1MB sequentially from memory	6 days <- do this!
Network round trip, same datacenter	11.5 days <- and this!
Read 1MB sequentially from disk	463 days



# Stateless




# Stateless



Memory we can use

# Stateless



Memory we can use

Mostly vertical  
scaling

# Why stateless?

# Why stateless?

- Workload (HTTP)

# Why stateless?

- Workload (HTTP)
- Hard to reuse memory

# Why stateless?

- Workload (HTTP)
- Hard to reuse memory
  - Short lived programs

# Why stateless?

- Workload (HTTP)
- Hard to reuse memory
  - Short lived programs
  - Single threaded programs



# Why stateless?

- Workload (HTTP)
- Hard to reuse memory
  - Short lived programs
  - Single threaded programs
- Tricky to coordinate servers

And for Elixir/Phoenix?

# And for Elixir/Phoenix?

- Workload (HTTP)

# And for Elixir/Phoenix?

- Workload (HTTP) —> Channels!

# And for Elixir/Phoenix?

- Workload (HTTP) —> Channels!
- Hard to reuse memory

# And for Elixir/Phoenix?

- Workload (HTTP) —> Channels!
- Hard to reuse memory
  - Short lived programs

# And for Elixir/Phoenix?

- Workload (HTTP) —> Channels!
- Hard to reuse memory
  - Short lived programs —> BEAM!

# And for Elixir/Phoenix?

- Workload (HTTP) —> Channels!
- Hard to reuse memory
  - Short lived programs —> BEAM!
  - Single threaded programs



# And for Elixir/Phoenix?

- Workload (HTTP) —> Channels!
- Hard to reuse memory
  - Short lived programs —> BEAM!
  - Single threaded programs —> BEAM!

# And for Elixir/Phoenix?



- Workload (HTTP) —> Channels!
- Hard to reuse memory
  - Short lived programs —> BEAM!
  - Single threaded programs —> BEAM!
- Tricky to coordinate servers

# And for Elixir/Phoenix?


- Workload (HTTP) —> Channels!
- Hard to reuse memory
  - Short lived programs —> BEAM!
  - Single threaded programs —> BEAM!
  - Tricky to coordinate servers —> BEAM!

Awesome! Let's build  
a stateful web app!



# A stateful web app



# A stateful web app





# A stateful web app



Data for user #42

# A stateful web app

user #42





Data for user #42



# A stateful web app



user #42  ?



Data for user #42


# A stateful web app

user #42 




Data for user #42

# A stateful web app



Data for user #42

# A stateful web app




```
send remote_pid,  
  {:get_user_data, 42}
```

Data for user #42

# A stateful web app

user #42  `GenServer.call(..., ...)`




```
send remote_pid,  
{:get_user_data, 42}
```

Data for user #42



**ARE YOU SURE?**





# DISTRIBUTED SYSTEMS





# Buzzword Bingo

Stateful

Distributed

Fault-tolerant

Real-time

Impress your cat

(application)

# Buzzword Bingo

Stateful - ✓

Distributed

Fault-tolerant

Real-time

Impress your cat

(application)

# Distributed Primitives and Patterns

# Distributed Primitives and Patterns

`send(remote_pid, ...), GenServer.call/cast`

distributed systems

# Distributed Primitives and Patterns

`send(remote_pid, ...), GenServer.call/cast`

distributed systems

$\approx$

# Distributed Primitives and Patterns

`send(remote_pid, ...), GenServer.call/cast`

distributed systems

$\approx$

`print("<div>")`

web development

# Distributed Primitives and Patterns

distributed systems

$\approx$

# Distributed Primitives and Patterns

ad hoc message passing or RPC

distributed systems

$\approx$



# Distributed Primitives and Patterns

ad hoc message passing or RPC

distributed systems

$\approx$

<?php

# Distributed Primitives and Patterns

ad hoc message passing or RPC

distributed systems

$\approx$

```
<?php  
    $foo = mysql_query($my_cool_query);
```

# Distributed Primitives and Patterns

ad hoc message passing or RPC

distributed systems

≈

```
<?php
    $foo = mysql_query($my_cool_query);
    echo "<div>$foo</div>";
```

# Distributed Primitives and Patterns

ad hoc message passing or RPC

distributed systems

≈

```
<?php
    $foo = mysql_query($my_cool_query);
    echo "<div>$foo</div>";
?>
```

# Distributed Primitives and Patterns

ad hoc message passing or RPC

distributed systems

≈

```
<?php
    $foo = mysql_query($my_cool_query);
    echo "<div>$foo</div>";
?>
```

web development

# Distributed Primitives and Patterns

```
<?php
    $foo = mysql_query($my_cool_query);
    echo "<div>$foo</div>";
?>
```

# Distributed Primitives and Patterns

\$\$\$\$\$\$\$\$

```
<?php
    $foo = mysql_query($my_cool_query);
    echo "<div>$foo</div>";
?>
```

# Distributed Primitives and Patterns

I want my MVC



# Distributed Primitives and Patterns

*I want my MVC*

what is the MVC of distributed systems?

# Distributed CAP

# Distributed CAP

Consistency

# Distributed CAP

Consistency

Availability

# Distributed CAP

Consistency

Availability

Partition tolerance

# Distributed CAP


Consistency

"Pick two"

Availability

Partition tolerance

# Distributed CAP



# Distributed CAP

CP

AP

CA



# Distributed CAP


CP

AP

CA

# Distributed CAP

CP



CA

# Distributed CAP

AP

# Distributed CAP

AP

gossip protocols

Distributed  
CAP

AP

gossip protocols

CRDTs

# Distributed CAP

AP

gossip protocols

CRDTs

distributed hash tables

# Distributed CAP

AP

gossip protocols

CRDTs

distributed hash tables

Awesome! Let's build a  
stateful, distributed web app  
using a distributed hash table!



A stateful, distributed web  
app with a DHT

A stateful, distributed web  
app with a DHT

Use a framework, don't write one

# Riak Core

# Riak Core

"a toolkit for building distributed, scalable, fault-tolerant applications"

- [riak\\_core README](#)

# Riak Core

mind blowing, advanced technology

# Riak Core

mind blowing, advanced technology

without peer in other platforms

# Riak Core

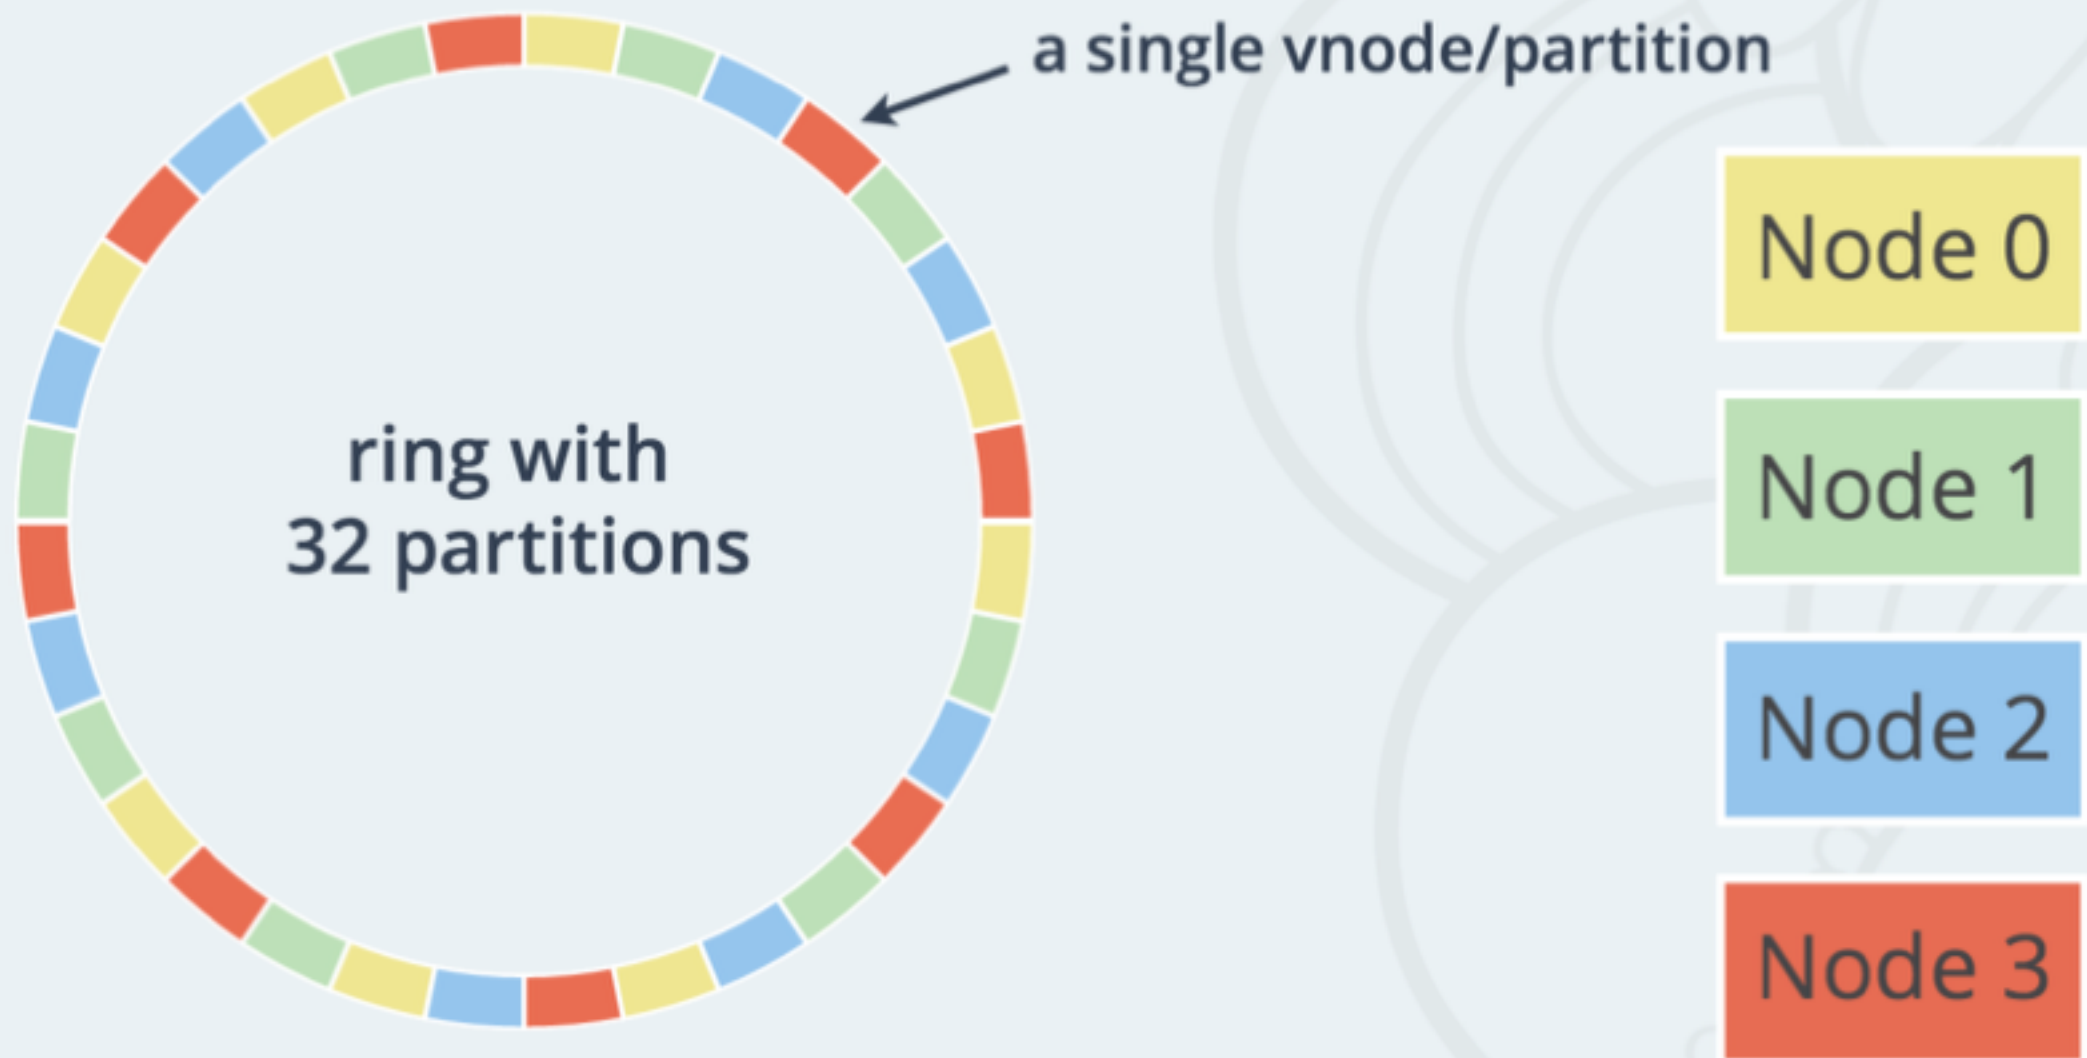
mind blowing, advanced technology

without peer in other platforms

source of magic for Basho + others

# Riak Core


## THE RIAK "RING" ARCHITECTURE






# Riak Core Hash Ring

```
my_hash = {}
```



# Riak Core Hash Ring


```
my_hash = {}
```



```
my_hash["answer"] = 42
```

# Riak Core Hash Ring


```
my_hash = {}
```



```
my_hash["answer"] = 42  
hash("answer") → 10403
```

# Riak Core Hash Ring


```
my_hash = {}
```



```
my_hash["answer"] = 42  
hash("answer") -> 10403  
index = 10403 % 4
```

# Riak Core Hash Ring

```
my_hash = {}
```



```
my_hash["answer"] = 42  
hash("answer") -> 10403  
index = 10403 % 4  
# put 42 in index 3!
```

# Riak Core Hash Ring

```
my_hash = {"answer"=>42}
```

|  |  |  |                 |
|--|--|--|-----------------|
|  |  |  | "answer",<br>42 |
|--|--|--|-----------------|

```
my_hash["answer"] = 42  
hash("answer") -> 10403  
index = 10403 % 4  
# put 42 in index 3!
```

# Riak Core Hash Ring

```
my_hash = {"answer"=>42}
```

|  |  |  |                 |
|--|--|--|-----------------|
|  |  |  | "answer",<br>42 |
|--|--|--|-----------------|

Which server owns which bucket?

# Riak Core Hash Ring

```
my_hash = {"answer"=>42}
```

|  |  |  |                 |
|--|--|--|-----------------|
|  |  |  | "answer",<br>42 |
|--|--|--|-----------------|

A

Which server owns which bucket?



# Riak Core Hash Ring

```
my_hash = {"answer"=>42}
```

|   |   |  |                 |
|---|---|--|-----------------|
|   |   |  | "answer",<br>42 |
| A | B |  |                 |

Which server owns which bucket?

# Riak Core Hash Ring

```
my_hash = {"answer"=>42}
```

|   |   |   |                 |
|---|---|---|-----------------|
|   |   |   | "answer",<br>42 |
| A | B | A |                 |

Which server owns which bucket?

# Riak Core Hash Ring


```
my_hash = {"answer"=>42}
```

|   |   |   |                 |
|---|---|---|-----------------|
|   |   |   | "answer",<br>42 |
| A | B | A | B               |

Which server owns which bucket?

# Riak Core Hash Ring


```
my_hash = {"answer"=>42}
```



Which server owns which bucket?

# Riak Core

## THE RIAK "RING" ARCHITECTURE



# A Riak Core App: pingring - the app

# A Riak Core App: pingring - the app

ping -> pong

# A Riak Core App: pingring - the app

ping -> pong

hash(timestamp) -> vnode (bucket)



# A Riak Core App: pingring - the app

ping -> pong

hash(timestamp) -> vnode (bucket)

example of distributing CPU work

# A Riak Core App: pingring - the parts

# A Riak Core App: pingring - the parts

service (high level API)

# A Riak Core App: pingring - the parts

service (high level API)

vnode (business logic)

# A Riak Core App: pingring - the parts

service (high level API)

vnode (business logic)

(supervisor)

# A Riak Core App: pingring - the parts

service (high level API)

vnode (business logic)

(supervisor)

(application)

# A Riak Core App: pingring - the parts

service (high level API)

vnode (business logic)

(supervisor)

(application)

# A Riak Core App: pingring - service

```
defmodule Pingring.Service do  
  def ping do
```

```
    end  
  end
```



# A Riak Core App: pingring - service

```
defmodule Pingring.Service do
  def ping do
    doc_idx = hash_key(
    end
  end
end
```

# A Riak Core App: pingring - service

```
defmodule Pingring.Service do
  def ping do
    doc_idx = hash_key(
      :os.timestamp )

    end
  end
end
```



# A Riak Core App: pingring - service

```
defmodule Pingring.Service do
  def ping do
    doc_idx = :riak_core_util.chash_key(
      {"ping", :erlang.term_to_binary(:os.timestamp)})

    end
  end
end
```

# A Riak Core App: pingring - service

```
defmodule Pingring.Service do
  def ping do
    doc_idx = :riak_core_util.chash_key(
      {"ping", :erlang.term_to_binary(:os.timestamp)})

    pref_list =

  end
end
```

# A Riak Core App: pingring - service

```
defmodule Pingring.Service do
  def ping do
    doc_idx = :riak_core_util.chash_key(
      {"ping", :erlang.term_to_binary(:os.timestamp)})

    pref_list = get_primary_apl(
      doc_idx,
    )

  end
end
```

# A Riak Core App: pingring - service

```
defmodule Pingring.Service do
  def ping do
    doc_idx = :riak_core_util.chash_key(
      {"ping", :erlang.term_to_binary(:os.timestamp)})

    pref_list = get_primary_apl(
      doc_idx, Pingring.Service)

  end
end
```

# A Riak Core App: pingring - service

```
defmodule Pingring.Service do
  def ping do
    doc_idx = :riak_core_util.chash_key(
      {"ping", :erlang.term_to_binary(:os.timestamp)})

    pref_list = get_primary_apl(
      doc_idx, 1, Pingring.Service)

  end
end
```



# A Riak Core App: pingring - service

```
defmodule Pingring.Service do
  def ping do
    doc_idx = :riak_core_util.chash_key(
      {"ping", :erlang.term_to_binary(:os.timestamp)})

    pref_list = :riak_core_apl.get_primary_apl(
      doc_idx, 1, Pingring.Service)

    end
  end
end
```

# A Riak Core App: pingring - service

```
defmodule Pingring.Service do
  def ping do
    doc_idx = :riak_core_util.chash_key(
      {"ping", :erlang.term_to_binary(:os.timestamp)})

    pref_list = :riak_core_apl.get_primary_apl(
      doc_idx, 1, Pingring.Service)

    [{index_node, _type}] = pref_list

  end
end
```

# A Riak Core App: pingring - service

```
defmodule Pingring.Service do
  def ping do
    doc_idx = :riak_core_util.chash_key(
      {"ping", :erlang.term_to_binary(:os.timestamp)})

    pref_list = :riak_core_apl.get_primary_apl(
      doc_idx, 1, Pingring.Service)

    [{index_node, _type}] = pref_list

    sync_spawn_command(
      index_node,

    end
  end
end
```

# A Riak Core App: pingring - service

```
defmodule Pingring.Service do
  def ping do
    doc_idx = :riak_core_util.chash_key(
      {"ping", :erlang.term_to_binary(:os.timestamp)})

    pref_list = :riak_core_apl.get_primary_apl(
      doc_idx, 1, Pingring.Service)

    [{index_node, _type}] = pref_list

    sync_spawn_command(
      index_node, :ping,
    )
  end
end
```

# A Riak Core App: pingring - service

```
defmodule Pingring.Service do
  def ping do
    doc_idx = :riak_core_util.chash_key(
      {"ping", :erlang.term_to_binary(:os.timestamp)})

    pref_list = :riak_core_apl.get_primary_apl(
      doc_idx, 1, Pingring.Service)

    [{index_node, _type}] = pref_list

    sync_spawn_command(
      index_node, :ping, Pingring.Vnode_master)
  end
end
```

# A Riak Core App: pingring - service

```
defmodule Pingring.Service do
  def ping do
    doc_idx = :riak_core_util.chash_key(
      {"ping", :erlang.term_to_binary(:os.timestamp)})

    pref_list = :riak_core_apl.get_primary_apl(
      doc_idx, 1, Pingring.Service)

    [{index_node, _type}] = pref_list
    # riak core appends "_master" to Pingring.Vnode.
    sync_spawn_command(
      index_node, :ping, Pingring.Vnode_master)
  end
end
```

# A Riak Core App: pingring - service

```
defmodule Pingring.Service do
  def ping do
    doc_idx = :riak_core_util.chash_key(
      {"ping", :erlang.term_to_binary(:os.timestamp)})

    pref_list = :riak_core_apl.get_primary_apl(
      doc_idx, 1, Pingring.Service)

    [{index_node, _type}] = pref_list
    # riak core appends "_master" to Pingring.Vnode.
    :riak_core_vnode_master.sync_spawn_command(
      index_node, :ping, Pingring.Vnode_master)
  end
end
```

# A Riak Core App: pingring - service

```
defmodule Pingring.Service do
  def ping do
    doc_idx = :riak_core_util.chash_key(
      {"ping", :erlang.term_to_binary(:os.timestamp)})

    pref_list = :riak_core_apl.get_primary_apl(
      doc_idx, 1, Pingring.Service)

    [{index_node, _type}] = pref_list
    # riak core appends "_master" to Pingring.Vnode.
    :riak_core_vnode_master.sync_spawn_command(
      index_node, :ping, Pingring.Vnode_master)
  end
end
```



# A Riak Core App: pingring - the parts

service (high level API)

vnode (business logic)

(supervisor)

(application)

# A Riak Core App: pingring - vnode

```
defmodule Pingring.Vnode do
```

```
end
```

# A Riak Core App: pingring - vnode

```
defmodule Pingring.Vnode do  
  @behaviour :riak_core_vnode
```

```
end
```

# A Riak Core App: pingring - vnode

```
defmodule Pingring.Vnode do  
  @behaviour :riak_core_vnode  
  # ... some boilerplate for startup
```

```
end
```

# A Riak Core App: pingring - vnode

```
defmodule Pingring.Vnode do
  @behaviour :riak_core_vnode
  # ... some boilerplate for startup
  def init(          ), do: {:ok, %{          }}

end
```

# A Riak Core App: pingring - vnode

```
defmodule Pingring.Vnode do
  @behaviour :riak_core_vnode
  # ... some boilerplate for startup
  def init([part]), do: {:ok, %{part: part}}

end
```

# A Riak Core App: pingring - vnode

```
defmodule Pingring.Vnode do
  @behaviour :riak_core_vnode
  # ... some boilerplate for startup
  def init([part]), do: {:ok, %{part: part}}

  def handle_command(

  ) do

  end

end
```

# A Riak Core App: pingring - vnode

```
defmodule Pingring.Vnode do
  @behaviour :riak_core_vnode
  # ... some boilerplate for startup
  def init([part]), do: {:ok, %{part: part}}

  def handle_command(
    :ping, _sender, %{part: part} = state
  ) do

  end

end
```



# A Riak Core App: pingring - vnode

```
defmodule Pingring.Vnode do
  @behaviour :riak_core_vnode
  # ... some boilerplate for startup
  def init([part]), do: {:ok, %{part: part}}

  def handle_command(
    :ping, _sender, %{part: part} = state
  ) do
    {:reply, {:pong, part}, state}
  end
end
```

end

# A Riak Core App: pingring - vnode

```
defmodule Pingring.Vnode do
  @behaviour :riak_core_vnode
  # ... some boilerplate for startup
  def init([part]), do: {:ok, %{part: part}}

  def handle_command(
    :ping, _sender, %{part: part} = state
  ) do
    {:reply, {:pong, part}, state}
  end
  # ... other callbacks for :riak_core_vnode
end
```

Demo

```
iex(dev_a@127.0.0.1)27> █
```

⌘

---

```
iex(dev_b@127.0.0.1)8>
```

---

What about...state?

# A Riak Core App: store\_fetch - the app

# A Riak Core App: store\_fetch - the app

```
store(key, data)
```

# A Riak Core App: store\_fetch - the app

store(key, data)

store in ETS



# A Riak Core App: store\_fetch - the app

store(key, data)

store in ETS

hash(key) -> vnode

# A Riak Core App: store\_fetch - service

```
def store(key, data) do
```

```
end
```

# A Riak Core App: store\_fetch - service

```
def store(key, data) do
  doc_idx = hash_key(
end
```

# A Riak Core App: store\_fetch - service

```
def store(key, data) do
  doc_idx =             hash_key(
                           key   )

end
```

# A Riak Core App: store\_fetch - service

```
def store(key, data) do
  doc_idx = :riak_core_util.chash_key(
    key
  )

end
```

# A Riak Core App: store\_fetch - service

```
def store(key, data) do
  doc_idx = :riak_core_util.chash_key(
    {"store", :erlang.term_to_binary(key)})

  end
```

# A Riak Core App: store\_fetch - service

```
def store(key, data) do
  doc_idx = :riak_core_util.chash_key(
    {"store", :erlang.term_to_binary(key)})

  pref_list = get_primary_apl(
    doc_idx,
    )

end
```

# A Riak Core App: store\_fetch - service

```
def store(key, data) do
  doc_idx = :riak_core_util.chash_key(
    {"store", :erlang.term_to_binary(key)})

  pref_list = get_primary_apl(
    doc_idx, StoreFetch.Service)

end
```



# A Riak Core App: store\_fetch - service

```
def store(key, data) do
  doc_idx = :riak_core_util.chash_key(
    {"store", :erlang.term_to_binary(key)})

  pref_list = get_primary_apl(
    doc_idx, 1, StoreFetch.Service)

end
```

# A Riak Core App: store\_fetch - service

```
def store(key, data) do
  doc_idx = :riak_core_util.chash_key(
    {"store", :erlang.term_to_binary(key)})

  pref_list = :riak_core_apl.get_primary_apl(
    doc_idx, 1, StoreFetch.Service)

end
```

# A Riak Core App: store\_fetch - service

```
def store(key, data) do
  doc_idx = :riak_core_util.chash_key(
    {"store", :erlang.term_to_binary(key)})

  pref_list = :riak_core_apl.get_primary_apl(
    doc_idx, 1, StoreFetch.Service)

  [{index_node, _type}] = pref_list

end
```

# A Riak Core App: store\_fetch - service

```
def store(key, data) do
  doc_idx = :riak_core_util.chash_key(
    {"store", :erlang.term_to_binary(key)})

  pref_list = :riak_core_apl.get_primary_apl(
    doc_idx, 1, StoreFetch.Service)

  [{index_node, _type}] = pref_list
    sync_spawn_command(
      index_node,
    )
end
```

# A Riak Core App: store\_fetch - service

```
def store(key, data) do
  doc_idx = :riak_core_util.chash_key(
    {"store", :erlang.term_to_binary(key)})

  pref_list = :riak_core_apl.get_primary_apl(
    doc_idx, 1, StoreFetch.Service)

  [{index_node, _type}] = pref_list
    sync_spawn_command(
      index_node, {:store, key, data},
    )
end
```

# A Riak Core App: store\_fetch - service

```
def store(key, data) do
  doc_idx = :riak_core_util.chash_key(
    {"store", :erlang.term_to_binary(key)})

  pref_list = :riak_core_apl.get_primary_apl(
    doc_idx, 1, StoreFetch.Service)

  [{index_node, _type}] = pref_list
    sync_spawn_command(
      index_node, {:store, key, data},
      StoreFetch.Vnode_master)
end
```

# A Riak Core App: store\_fetch - service

```
def store(key, data) do
  doc_idx = :riak_core_util.chash_key(
    {"store", :erlang.term_to_binary(key)})

  pref_list = :riak_core_apl.get_primary_apl(
    doc_idx, 1, StoreFetch.Service)

  [{index_node, _type}] = pref_list
  :riak_core_vnode_master.sync_spawn_command(
    index_node, {:store, key, data},
    StoreFetch.Vnode_master)
end
```

# A Riak Core App: store\_fetch - vnode

```
defmodule StoreFetch.Vnode do
```

```
end
```



# A Riak Core App: store\_fetch - vnode

```
defmodule StoreFetch.Vnode do  
  @behaviour :riak_core_vnode  
  # ... some boilerplate for startup
```

```
end
```

# A Riak Core App: store\_fetch - vnode

```
defmodule StoreFetch.Vnode do
  @behaviour :riak_core_vnode
  # ... some boilerplate for startup
  def init([_part]) do
```

```
end
```

```
end
```

# A Riak Core App: store\_fetch - vnode

```
defmodule StoreFetch.Vnode do
  @behaviour :riak_core_vnode
  # ... some boilerplate for startup
  def init([_part]) do
    ets_handle = :ets.new(nil, [])
    {:ok, %{db: ets_handle}}
  end
end
```

end

# A Riak Core App: store\_fetch - vnode

```
defmodule StoreFetch.Vnode do
  @behaviour :riak_core_vnode
  # ... some boilerplate for startup
  def init([_part]) do
    ets_handle = :ets.new(nil, [])
    {:ok, %{db: ets_handle}}
  end
  def handle_command(

  ) do

  end

end
```

# A Riak Core App: store\_fetch - vnode

```
defmodule StoreFetch.Vnode do
  @behaviour :riak_core_vnode
  # ... some boilerplate for startup
  def init([_part]) do
    ets_handle = :ets.new(nil, [])
    {:ok, %{db: ets_handle}}
  end
  def handle_command(
    {:store, key, data}, _sender, %{db: db} = state
  ) do

  end
end
```

# A Riak Core App: store\_fetch - vnode


```
defmodule StoreFetch.Vnode do
  @behaviour :riak_core_vnode
  # ... some boilerplate for startup
  def init([_part]) do
    ets_handle = :ets.new(nil, [])
    {:ok, %{db: ets_handle}}
  end
  def handle_command(
    {:store, key, data}, _sender, %{db: db} = state
  ) do
    result = :ets.insert(db, {key, data})
    {:reply, result, state}
  end
end
```

# A Riak Core App: store\_fetch - vnode

```
defmodule StoreFetch.Vnode do
  @behaviour :riak_core_vnode
  # ... some boilerplate for startup
  def init([_part]) do
    ets_handle = :ets.new(nil, [])
    {:ok, %{db: ets_handle}}
  end
  def handle_command(
    {:store, key, data}, _sender, %{db: db} = state
  ) do
    result = :ets.insert(db, {key, data})
    {:reply, result, state}
  end
  # ... same for fetch, but :ets.lookup instead
end
```

**I HAVE**


**MY DOUBTS**





**NINE NINES?**

**MORE LIKE...  
NOM NOMS**





# Buzzword Bingo

Stateful - ✓

Distributed

Fault-tolerant

Real-time

Impress your cat

(application)

# Buzzword Bingo

Stateful - ✓

Distributed - ✓

Fault-tolerant

Real-time

Impress your cat

(application)

# Fault Tolerance

Computers needed for fault tolerance?

# Fault Tolerance

Computers needed for fault tolerance?

$> 1$

# A Riak Core App: store\_fetch - the parts

service

vnode

(supervisor)

(application)

# A Riak Core App: store\_fetch - the parts

service

vnode

(supervisor)

(application)

write coordinator (plus supervisor)

# The Write Coordinator



# The Write Coordinator

executes commands on multiple vnodes

# The Write Coordinator

executes commands on multiple vnodes

riak\_core takes care of spreading vnodes across  
servers

# The Write Coordinator

executes commands on multiple vnodes

riak\_core takes care of spreading vnodes across  
servers

(as much as possible)

# Using the Write Coordinator

```
defmodule StoreFetch.Service do
```

```
end
```

# Using the Write Coordinator

```
defmodule StoreFetch.Service do
  def store(key, data) do
```

```
    end
```

```
end
```

# Using the Write Coordinator

```
defmodule StoreFetch.Service do  
  def store(key, data, n, w) do
```

```
    end
```

```
  end
```

# Using the Write Coordinator

```
defmodule StoreFetch.Service do
  def store(key, data, n, w) do
    {:ok, req_id} = StoreFetch.WCoord.do(
      )

  end
end
```

# Using the Write Coordinator

```
defmodule StoreFetch.Service do
  def store(key, data, n, w) do
    {:ok, req_id} = StoreFetch.WCoord.do(
      key, {:store, key, data}, n, w)

    end
end
```



# Using the Write Coordinator

```
defmodule StoreFetch.Service do
  def store(key, data, n, w) do
    {:ok, req_id} = StoreFetch.WCoord.do(
      key, {:store, key, data}, n, w)

    receive do

    after

    end
  end
end
```

# Using the Write Coordinator

```
defmodule StoreFetch.Service do
  def store(key, data, n, w) do
    {:ok, req_id} = StoreFetch.WCoord.do(
      key, {:store, key, data}, n, w)

    receive do

      after
        5000 ->
          {:error, :timeout}
      end
    end
  end
end
```

# Using the Write Coordinator

```
defmodule StoreFetch.Service do
  def store(key, data, n, w) do
    {:ok, req_id} = StoreFetch.WCoord.do(
      key, {:store, key, data}, n, w)

    receive do
      {^req_id, value} ->
        {:ok, value}
    after
      5000 ->
        {:error, :timeout}
    end
  end
end
```

# Using the Write Coordinator

```
defmodule StoreFetch.Service do
  def store(key, data, n, w) do
    {:ok, req_id} = StoreFetch.WCoord.do(
      key, {:store, key, data}, n, w)

    receive do
      {^req_id, value} ->
        {:ok, value}
    after
      5000 ->
        {:error, :timeout}
    end
  end
end

# ... fetch implementation
end
```

# The Write Coordinator

# The Write Coordinator

:gen\_fsm

# The Write Coordinator

`:gen_fsm`

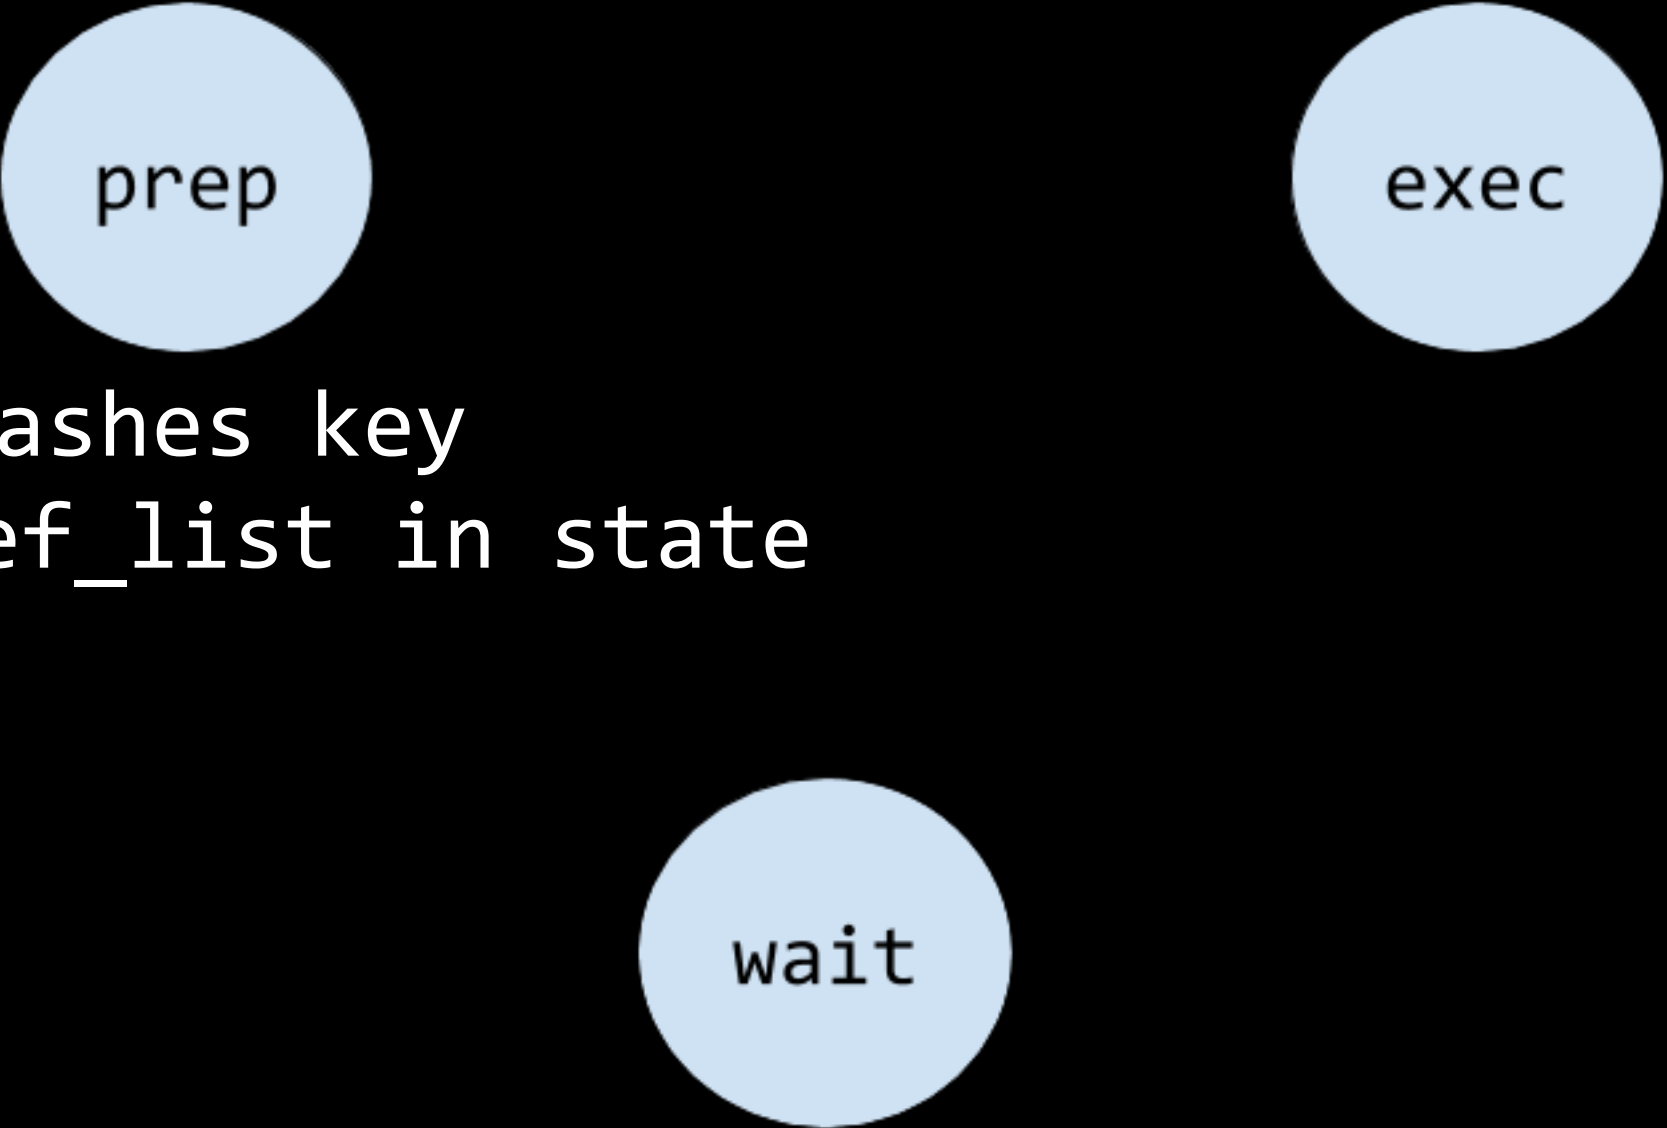
spawned on demand (`:simple_one_for_one`)

# The Write Coordinator





# The Write Coordinator



```
graph TD; prep((prep)) --> exec((exec)); exec --> wait((wait)); wait --> prep;
```

prep

exec

hashes key  
puts pref\_list in state

wait

# The Write Coordinator



hashes key  
puts pref\_list in state



# The Write Coordinator

sends command to pref\_list




hashes key  
puts pref\_list in state



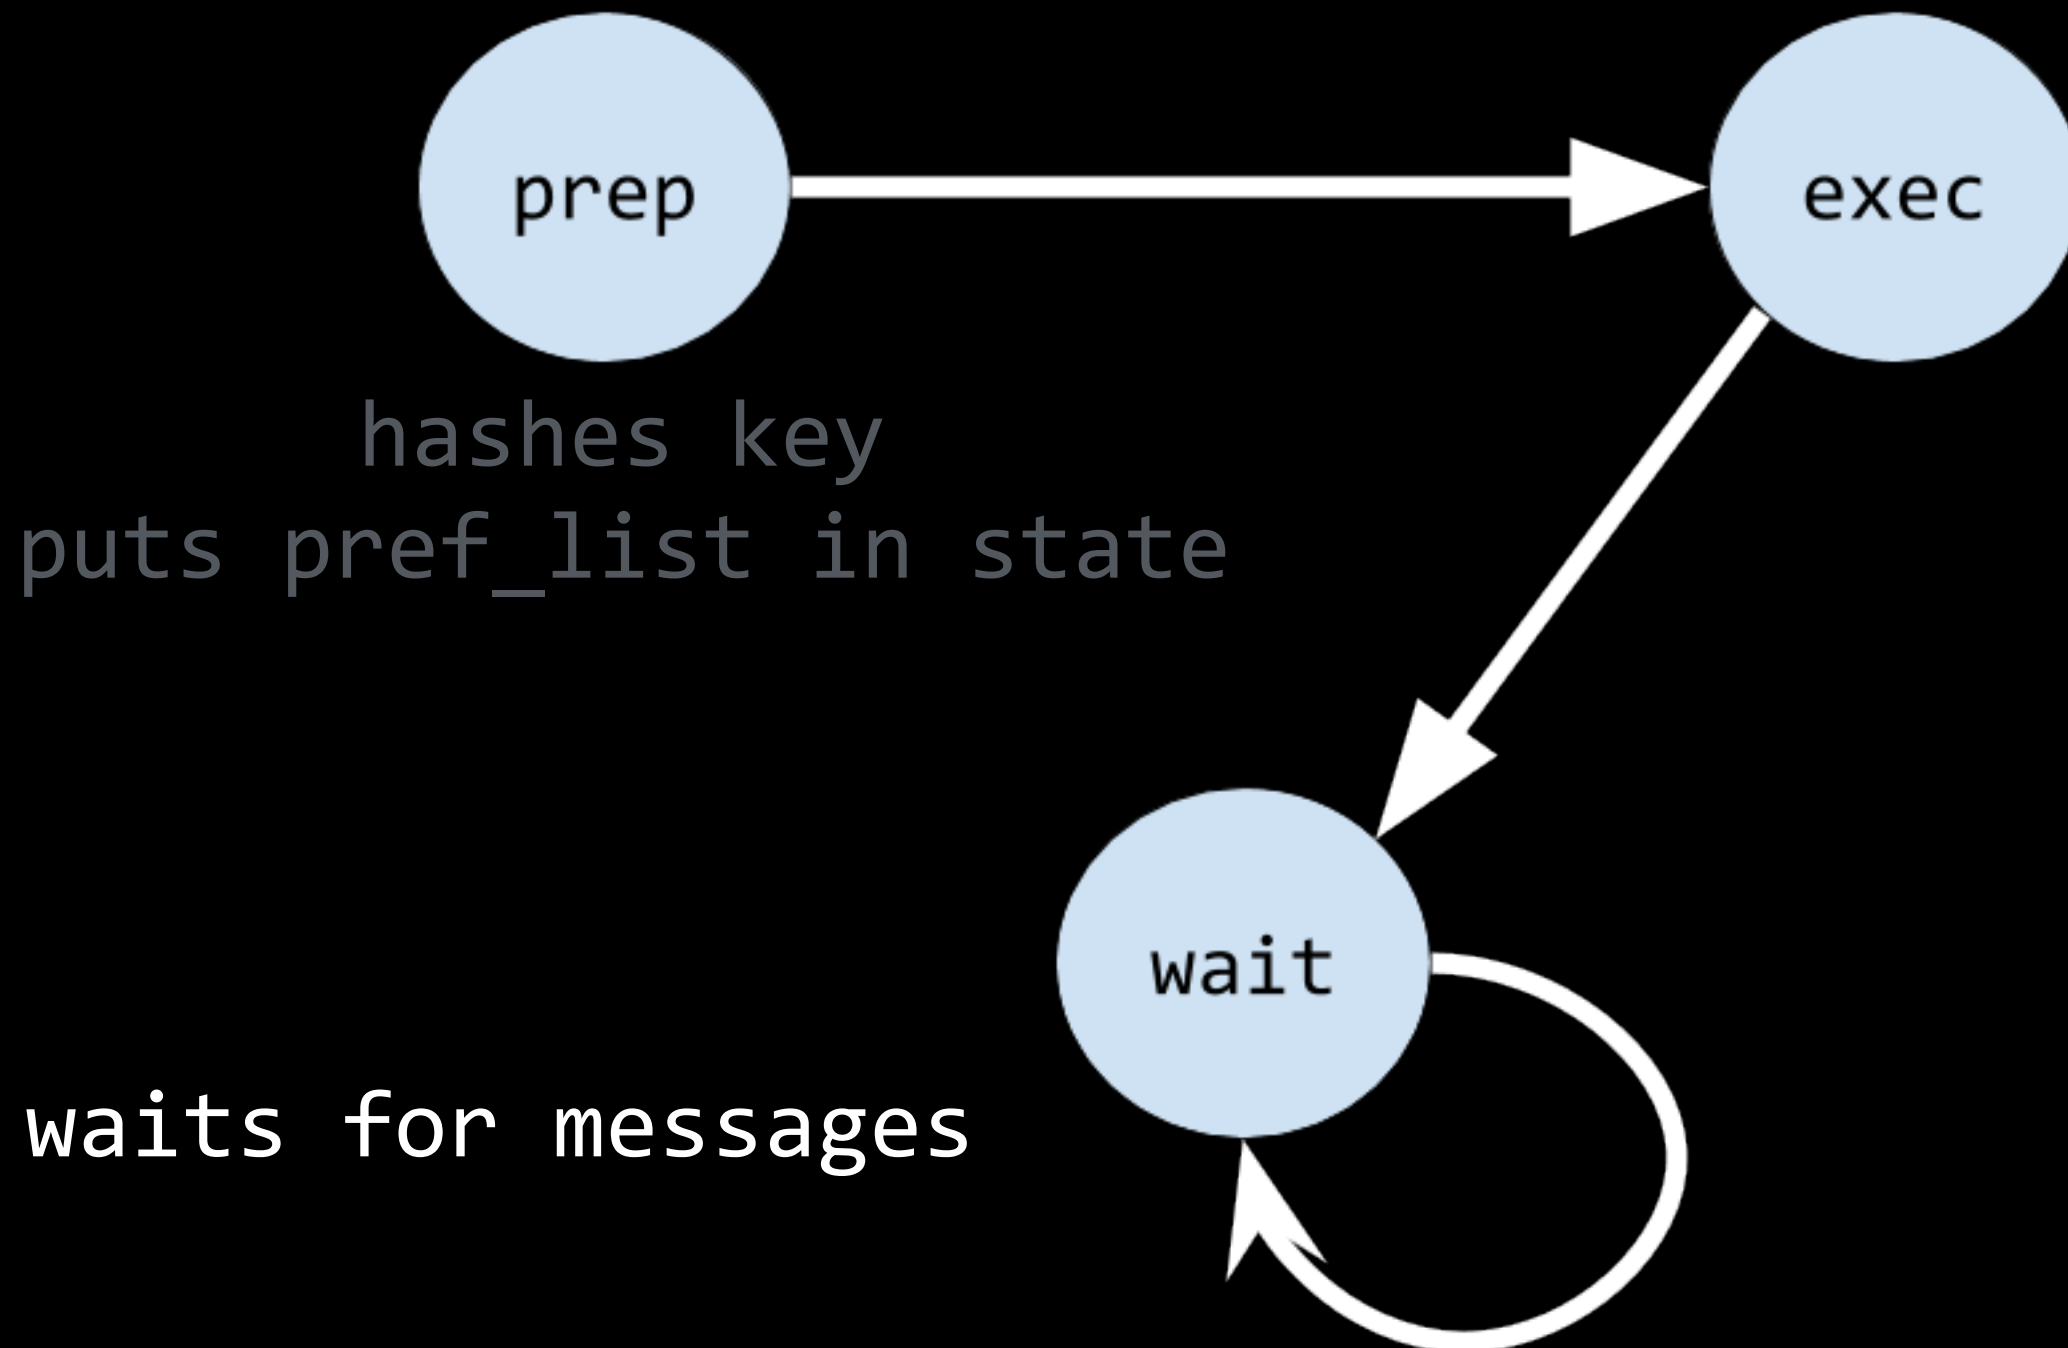
# The Write Coordinator

sends command to pref\_list



# The Write Coordinator

sends command to pref\_list




# The Write Coordinator

sends command to pref\_list



hashes key  
puts pref\_list in state



waits for messages  
#responses > w? done!

# Fault Tolerance handoff

# Fault Tolerance handoff

what if we add (or lose) servers?



# Fault Tolerance handoff

what if we add (or lose) servers?

"handing off" responsibility for vnode

# Fault Tolerance handoff

what if we add (or lose) servers?

"handing off" responsibility for vnode

series of callbacks in Vnode module

# Fault Tolerance handoff

what if we add (or lose) servers?

"handing off" responsibility for vnode

series of callbacks in Vnode module

mostly a matter of serialisation

# Buzzword Bingo

Stateful - ✓

Distributed - ✓

Fault-tolerant

Real-time

Impress your cat

(application)

# Buzzword Bingo

Stateful - ✓

Distributed - ✓

Fault-tolerant - ✓

Phoenix!

Real-time

Impress your cat

(application)

# Phoenix + Riak Core

# Phoenix + Riak Core

just use an umbrella!

# Phoenix + Riak Core

just use an umbrella!

then use the Service API in your Phoenix app  
somewhere



# Phoenix + Riak Core

just use an umbrella!

then use the Service API in your Phoenix app  
somewhere

```
scope "/api", MyApp do
  pipe_through :api
  put "/store/:key", StoreController, :store
  get "/store/:key", StoreController, :fetch
end
```

# Phoenix + Riak Core

```
curl -XPUT -d '{"a":"b"}' localhost:4000/api/store/my_key
```

# Phoenix + Riak Core

```
curl -XPUT -d '{"a":"b"}' localhost:4000/api/store/my_key
```

```
defmodule MyApp.StoreController do  
  use Phoenix.Controller  
  use MyApp.Web, :controller
```

```
end
```

# Phoenix + Riak Core

```
curl -XPUT -d '{"a":"b"}' localhost:4000/api/store/my_key
```

```
defmodule MyApp.StoreController do
```

```
  use Phoenix.Controller
```

```
  use MyApp.Web, :controller
```

```
  def store(
```

```
    ) do
```

```
end
```

```
end
```

# Phoenix + Riak Core

```
curl -XPUT -d '{"a":"b"}' localhost:4000/api/store/my_key
```

```
defmodule MyApp.StoreController do
  use Phoenix.Controller
  use MyApp.Web, :controller
  def store(
    %Plug.Conn{body_params: data}=conn,

  ) do

  end

end
```

# Phoenix + Riak Core

```
curl -XPUT -d '{"a":"b"}' localhost:4000/api/store/my_key
```

```
defmodule MyApp.StoreController do
  use Phoenix.Controller
  use MyApp.Web, :controller
  def store(
    %Plug.Conn{body_params: data}=conn,
    %{"key" => key}=params
  ) do
```

```
end
```

```
end
```

# Phoenix + Riak Core

```
curl -XPUT -d '{"a":"b"}' localhost:4000/api/store/my_key
```

```
defmodule MyApp.StoreController do
  use Phoenix.Controller
  use MyApp.Web, :controller
  def store(
    %Plug.Conn{body_params: data}=conn,
    %{"key" => key}=params
  ) do
    n = 3

  end
end
```

# Phoenix + Riak Core

```
curl -XPUT -d '{"a":"b"}' localhost:4000/api/store/my_key
```

```
defmodule MyApp.StoreController do
```

```
  use Phoenix.Controller
```

```
  use MyApp.Web, :controller
```

```
  def store(  
    %Plug.Conn{body_params: data}=conn,  
    %{"key" => key}=params  
  ) do
```

```
    n = 3
```

```
    result = StoreFetch.Service.store(key, data, n)
```

```
  end
```

```
end
```



# Phoenix + Riak Core

```
curl -XPUT -d '{"a":"b"}' localhost:4000/api/store/my_key
```

```
defmodule MyApp.StoreController do
  use Phoenix.Controller
  use MyApp.Web, :controller
  def store(
    %Plug.Conn{body_params: data}=conn,
    %{"key" => key}=params
  ) do
    n = 3
    result = StoreFetch.Service.store(key, data, n)
    render conn, store: result
  end
end
```

# Phoenix + Riak Core

```
curl -XPUT -d '{"a":"b"}' localhost:4000/api/store/my_key
```

```
defmodule MyApp.StoreController do
  use Phoenix.Controller
  use MyApp.Web, :controller
  def store(
    %Plug.Conn{body_params: data}=conn,
    %{"key" => key}=params
  ) do
    n = 3
    result = StoreFetch.Service.store(key, data, n)
    render conn, store: result
  end
  # ... similar for fetch/2
end
```

Demo

\$ █

---

```
iex(dev_a@127.0.0.1)8>
```

---

```
iex(dev_b@127.0.0.1)7>
```

# Buzzword Bingo

Stateful - ✓

Distributed - ✓

Fault-tolerant - ✓

Phoenix!

Real-time

Impress your cat

(application)

# Buzzword Bingo

Stateful - ✓

Distributed - ✓

Fault-tolerant - ✓

Phoenix! - ✓

Real-time

Impress your cat

(application)

# Real-time

# Real-time

can we do something with channels?



# Real-time

can we do something with channels?

need something to hash on (a key)

# Real-time

can we do something with channels?

need something to hash on (a key)

```
%Phoenix.Socket.Broadcast{  
  event: "new_msg",  
  payload: %{body: "hey everyone!"},  
  topic: "rooms:lobby"  
}
```

# Real-time

can we do something with channels?

need something to hash on (a key)

```
%Phoenix.Socket.Broadcast{  
  event: "new_msg",  
  payload: %{body: "hey everyone!"},  
  topic: "rooms:lobby"  
}
```

Phoenix.PubSub.Ricor

# Phoenix.PubSub.Ricor

hack hack Phoenix.PubSub adapter

# Phoenix.PubSub.Ricor

hack hack Phoenix.PubSub adapter

uses a Riak Core Service to direct messages to  
vnode (by topic)

# Phoenix.PubSub.Ricor

hack hack Phoenix.PubSub adapter

uses a Riak Core Service to direct messages to  
vnode (by topic)

vnode manages subscriptions and broadcasts

# Phoenix.PubSub.Ricor

```
def subscribe(pid, topic, opts) do
```

```
end
```

```
def unsubscribe(pid, topic) do
```

```
end
```

```
def broadcast(pid, topic, message) do
```

```
end
```



# Phoenix.PubSub.Ricor

```
def subscribe(pid, topic, opts) do
  Pubring.Service.subscribe(pid, topic, opts)
end
```

```
def unsubscribe(pid, topic) do
  Pubring.Service.unsubscribe(pid, topic)
end
```

```
def broadcast(pid, topic, message) do
  Pubring.Service.broadcast(topic, message)
end
```

# Pubring.Service

just like other services:

# Pubring.Service

just like other services:

- 1) hash the key (topic)

# Pubring.Service

just like other services:

- 1) hash the key (topic)
- 2) get preference list

# Pubring.Service

just like other services:

- 1) hash the key (topic)
- 2) get preference list
- 3) send command to vnode from list

# Pubring.Service

just like other services:

- 1) hash the key (topic)
- 2) get preference list
- 3) send command to vnode from list

let's look at the vnode!

# Pubring.Service

# Pubring.Vnode

```
defmodule Pubring.Vnode do  
  @behaviour :riak_core_vnode
```

```
end
```



# Pubring.Vnode

```
defmodule Pubring.Vnode do
  @behaviour :riak_core_vnode
  def handle_command(

  ) do

  end
  def handle_command(

  ) do

  end
end
end
```

# Pubring.Vnode

```
defmodule Pubring.Vnode do
  @behaviour :riak_core_vnode
  def handle_command(
    :subscribe
  ) do

  end

  def handle_command(
    :broadcast
  ) do

  end

end

end
```

# Pubring.Vnode

```
defmodule Pubring.Vnode do
  @behaviour :riak_core_vnode
  def handle_command(
    {:subscribe, pid, topic}, _sender, %{db: db}=state
  ) do

  end

  def handle_command(
    :broadcast
  ) do

  end

end

end
```

# Pubring.Vnode

```
defmodule Pubring.Vnode do
  @behaviour :riak_core_vnode
  def handle_command(
    {:subscribe, pid, topic}, _sender, %{db: db}=state
  ) do
    result = :ets.insert(db, {topic, pid})

  end
  def handle_command(
    :broadcast
  ) do

  end
end
end
```

# Pubring.Vnode

```
defmodule Pubring.Vnode do
  @behaviour :riak_core_vnode
  def handle_command(
    {:subscribe, pid, topic}, _sender, %{db: db}=state
  ) do
    result = :ets.insert(db, {topic, pid})
    {:reply, result, state}
  end
  def handle_command(
    :broadcast
  ) do

  end
end
```

end

# Pubring.Vnode

```
defmodule Pubring.Vnode do
  @behaviour :riak_core_vnode
  def handle_command(
    {:subscribe, pid, topic}, _sender, %{db: db}=state
  ) do
    result = :ets.insert(db, {topic, pid})
    {:reply, result, state}
  end
  def handle_command(
    :broadcast
  ) do

  end
end
```

end

# Pubring.Vnode

```
defmodule Pubring.Vnode do
  @behaviour :riak_core_vnode
  def handle_command(
    {:subscribe, pid, topic}, _sender, %{db: db}=state
  ) do
    result = :ets.insert(db, {topic, pid})
    {:reply, result, state}
  end
  def handle_command(
    {:broadcast, topic, msg}, _sender, %{db: db}=state
  ) do

  end
end
```

# Pubring.Vnode

```
defmodule Pubring.Vnode do
  @behaviour :riak_core_vnode
  def handle_command(
    {:subscribe, pid, topic}, _sender, %{db: db}=state
  ) do
    result = :ets.insert(db, {topic, pid})
    {:reply, result, state}
  end
  def handle_command(
    {:broadcast, topic, msg}, _sender, %{db: db}=state
  ) do
    :ets.match(db, {topic, :"$1"})

  end
end
```



# Pubring.Vnode

```
defmodule Pubring.Vnode do
  @behaviour :riak_core_vnode
  def handle_command(
    {:subscribe, pid, topic}, _sender, %{db: db}=state
  ) do
    result = :ets.insert(db, {topic, pid})
    {:reply, result, state}
  end
  def handle_command(
    {:broadcast, topic, msg}, _sender, %{db: db}=state
  ) do
    for [pid] <- :ets.match(db, {topic, :"$1"}) do

    end

  end
end
end
```

# Pubring.Vnode

```
defmodule Pubring.Vnode do
  @behaviour :riak_core_vnode
  def handle_command(
    {:subscribe, pid, topic}, _sender, %{db: db}=state
  ) do
    result = :ets.insert(db, {topic, pid})
    {:reply, result, state}
  end
  def handle_command(
    {:broadcast, topic, msg}, _sender, %{db: db}=state
  ) do
    for [pid] <- :ets.match(db, {topic, :"$1"}) do
      send(pid, msg)
    end
  end
end
end
```

# Pubring.Vnode

```
defmodule Pubring.Vnode do
  @behaviour :riak_core_vnode
  def handle_command(
    {:subscribe, pid, topic}, _sender, %{db: db}=state
  ) do
    result = :ets.insert(db, {topic, pid})
    {:reply, result, state}
  end
  def handle_command(
    {:broadcast, topic, msg}, _sender, %{db: db}=state
  ) do
    for [pid] <- :ets.match(db, {topic, :"$1"}) do
      send(pid, msg)
    end
    {:reply, :ok, state}
  end
end

end
```

Demo

1. [0:6 ml905199] beam.smp (tmux)

iex(dev\_a@127.0.0.1)16> # client connected here

iex(dev\_b@127.0.0.1)18> # hosts vnode for "private 99"

iex(dev\_c@127.0.0.1)19> # hosts vnode "lobby"

[0] <3:template-with-ngZ 4:bashZ 5:bash 6:beam.smp\*> "ml905199" 01:54 10-May-16

Hello Coffeeshop! x  
localhost:4000

Lobby

Private 99

localhost:4001

Lobby

Private 99

Hello Coffeeshop! x  
localhost:4002

Lobby

Private 99

Phoenix.PubSub.Ricor

# Phoenix.PubSub.Ricor

so far, just a really complicated  
Phoenix.PubSub.PG2

# Phoenix.PubSub.Ricor

so far, just a really complicated  
Phoenix.PubSub.PG2

but...state?



# Phoenix.PubSub.Ricor

so far, just a really complicated  
Phoenix.PubSub.PG2

but...state?

message history, game state, etc.

# Phoenix.PubSub.Ricor

so far, just a really complicated  
Phoenix.PubSub.PG2

but...state?

message history, game state, etc.

**superpower!**

# Buzzword Bingo

Stateful - ✓

Distributed - ✓

Fault-tolerant - ✓

Phoenix! - ✓

Real-time

Impress your cat

(application)

# Buzzword Bingo

Stateful - ✓

Distributed - ✓

Fault-tolerant - ✓

Phoenix! - ✓

Real-time - ✓

Impress your cat

(application)

A close-up photograph of a black and white cat, possibly a Persian or similar breed, with a grumpy or disapproving expression. The cat has large, yellow-green eyes and a pink nose. It is looking directly at the camera. The background is dark and out of focus. The text "{:ERROR, :NOPE}" is overlaid in white, bold, sans-serif font across the lower part of the cat's face.

**{:ERROR, :NOPE}**

# Buzzword Bingo

Stateful - ✓

Distributed - ✓

Fault-tolerant - ✓

Phoenix! - ✓

Real-time - ✓

Impress your cat

(application)

# Buzzword Bingo

Stateful - ✓

Distributed - ✓

Fault-tolerant - ✓

Phoenix! - ✓

Real-time - ✓

Impress your cat - :(

(application)

# Future



# Future

```
mix ricor.new MyApp
```

# Future

```
mix ricor.new MyApp
```

```
--phoenix
```

# Future

```
mix ricor.new MyApp
```

```
--phoenix
```

```
--write-coord
```

# Future

```
mix ricor.new MyApp
```

```
--phoenix
```

```
--write-coord
```

## GenVnode

# Future

```
mix ricor.new MyApp
```

```
--phoenix
```

```
--write-coord
```

```
GenVnode
```

Phoenix.PubSub.Ricor.ButSerious

# Future

```
mix ricor.new MyApp
```

```
--phoenix
```

```
--write-coord
```

```
GenVnode
```

```
Phoenix.PubSub.Ricor.ButSerious
```

???

Future

(your app here)

# Thanks!

- Mariano Guerra - Little Riak Core book, rebar3 Ricor template, talks
- Mark Allen - Udon, Ricor talk, Basho blog
- Ryan Zezeski - 'Try Try Try' blog
- Project FIFO - a Riak Core that compiles on Erlang 18
- All the comments in the riak\_core source

`{:exit, :talk_over}`



# Talk Materials

<https://github.com/kanatohodets/scalable-stateful-web-phoenix-riak-core-talk> <- not quite a transcript

[https://github.com/kanatohodets/elixir\\_riak\\_core\\_ping](https://github.com/kanatohodets/elixir_riak_core_ping)

<https://github.com/kanatohodets/phoenix-ricor-kv>

<https://github.com/kanatohodets/hashpub>